

Preprocessing Review

Why do we **preprocess** data when we build machine learning pipelines?

We preprocess data for two principle reasons:

1. To transform the data to better suit a model's underlying assumptions.
2. To format the data in the way a model expects.

Today, we're concerned with this second reason.

Inputs to Neural Networks

What does the input to a neural network look like?

Inputs to neural networks are **vectors**. Each entry in the vector corresponds to a feature, which the net uses to make predictions.

Crucially, these vectors contain can contain only *numerical* data. They *cannot* contain string data.

One-Hot Encoding

This poses a problem when we want to train a neural network on categorical data as we have in `indeed_job_dataset`.

The class columns contain values below:

Queried_Salary:

```
array(['<80000', '80000-99999', '100000-119999', '120000-139999',  
      '140000-159999', '>160000'], dtype=object)
```

Job_Type:

```
array(['data_scientist', 'data_analyst', 'data_engineer'], dtype=object)
```

Skill:

```
array(["['SAP', 'SQL']",  
      "['Machine Learning', 'R', 'SAS', 'SQL', 'Python']",  
      "['Data Mining', 'Data Management', 'R', 'SAS', 'SQL', 'STATA',  
      'SPSS', 'Data Analysis', 'Python']",  
      ...,  
      "['Spring', 'Data Management', 'Hadoop', 'Kafka', 'Java',  
      'Cassandra', 'S3', 'Spark', 'Kubernetes', 'CI', 'Design Experience',  
      'NoSQL', 'AWS']",  
      "['Spring', 'Ruby', 'Test Automation', 'Scripting', 'SDLC',  
      'Scala', 'Perl', 'Kafka', 'Management Experience', 'Ansible', 'Java',  
      'Docker', 'Python', 'Military Experience', 'Jenkins']",  
      "['JavaScript', 'TS/SCI Clearance', 'XML', 'Hadoop', 'HTML5',  
      'JSON', 'Java', 'Software Development', 'NoSQL', 'AJAX', 'CSS']"],  
      dtype=object)
```

Location:

```
array(['MO', 'TX', 'OR', 'DC', 'MD', 'NY', 'GA', 'ID', 'PA', 'FL', 'MA',  
      'VA', 'NJ', 'LA', 'CA', 'MN', 'WA', 'NC', 'IL', 'CO', 'UT', 'OH',  
      'USA', 'AR', 'ME', 'NV', 'CT', 'REMOTE', 'RI', 'TN', 'WI', 'SC',  
      'MI', 'KY', 'AZ', 'NE', 'IN', 'NM', 'AL', 'KS', 'IA', 'DE', nan,  
      'HI', 'NH', 'OK', 'VT', 'WV', 'SD', 'WY', 'MT', 'ND'],  
      dtype=object)
```

As these are not numerical values, we can't use them to fit our neural network. To fix this, we must convert each class label to a numerical value.

We do this via the following steps:

1. **Label Encoding.** First, we convert the possible classes to integer labels. E.g., 'data_analyst' will be 1; 'data_engineer', 2; and 'data_scientist', 3.
2. **One-Hot Encoding.** Then, we set each row's class value to an *array*. This array will have a 1 in whichever slot corresponds to the integer label. E.g., after one-hot encoding, a row with the class iris-setosa will have the array [1, 0, 0]. A row with class iris-virginica, the array [0, 0, 1]; etc.

In many cases, categories in the data sets you work with will already be label-encoded like we have for some skills and locations label encoded columns. In this case, we applied one-hot encoding to all string columns.

LET'S START!

It is really important to scale our data before using multilayer perceptron models.

Without scaling, it is often difficult for the training cycle to converge.

```
labels=["Job_Type", "Job_Title", "Company"]  
X = df.drop(labels=labels, axis=1)  
y = df["Job_Type"]  
print(X.shape, y.shape)  
(5715, 925) (5715,)
```

1) Scaling

We use MinMaxScaler

```
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import LabelEncoder, MinMaxScaler  
from tensorflow.keras.utils import to_categorical
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, random_state=1)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
(4286, 925) (1429, 925) (4286,) (1429,)
X_scaler = MinMaxScaler().fit(X_train)
X_train_scaled = X_scaler.transform(X_train)
X_test_scaled = X_scaler.transform(X_test)
```

2) One-hot encode the labels

```
label_encoder = LabelEncoder()
label_encoder.fit(y_train)
encoded_y_train = label_encoder.transform(y_train)
encoded_y_test = label_encoder.transform(y_test)
```

3) Creating our Model

We must first decide what kind of model to apply to our data.

For numerical data, we use a regressor model.

For categorical data, we use a classifier model.

In our project, we use a classifier to build our network and

defining our Model Architecture (the layers), we first need to create a sequential model.

```
# Create model and add layers

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

model = Sequential()
```

Next, we add our first layer. This layer requires you to specify both the number of inputs and the number of nodes that you want in the hidden layer.

```
model.add(Dense(units=10, activation='relu', input_dim=X_features))
units → number of hidden nodes
x_features → number of our inputs
```

we add second layer

```
model.add(Dense(units=10, activation='relu'))
```

Our final layer is the output layer. Here, we need to specify the activation function (softmax is our classification) and the number of classes (labels) that we are trying to predict (y_features).

```
model.add(Dense(units=y_features, activation='softmax'))  
y_features → Job_Types
```

4) Compile the model

Now that we have our model architecture defined, we must compile the model using a loss function and optimizer. We can also specify additional training metrics such as accuracy.

We use categorical crossentropy for categorical data and mean squared error for regression
Hint: your output layer in this example is using software for logistic regression (categorical)

```
model.compile(optimizer='adam',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

5) Model Summary

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 10)	9260
dense_1 (Dense)	(None, 10)	110
dense_2 (Dense)	(None, 3)	33

```
Total params: 9,403  
Trainable params: 9,403  
Non-trainable params: 0
```

6) Training the model

Finally, we train our model using our training data

Training consists of updating our weights using our optimizer and loss function. In our project, we choose 20 iterations (loops) of training that are called epochs.

We also choose to shuffle our training data and increase the detail printed out during each training cycle.

```
model.fit(  
    X_train_scaled,  
    y_train_categorical,  
    epochs=20,  
    shuffle=True  
)
```

```
Epoch 1/20  
134/134 [=====] - 0s 1ms/step - loss: 0.9538 -  
accuracy: 0.5343  
Epoch 2/20  
134/134 [=====] - 0s 1ms/step - loss: 0.6043 -  
accuracy: 0.7928  
Epoch 3/20  
134/134 [=====] - 0s 1ms/step - loss: 0.4468 -  
accuracy: 0.8327  
Epoch 4/20  
134/134 [=====] - 0s 1ms/step - loss: 0.3782 -  
accuracy: 0.8560  
Epoch 5/20  
134/134 [=====] - 0s 1ms/step - loss: 0.3375 -  
accuracy: 0.8745  
Epoch 6/20  
134/134 [=====] - 0s 1ms/step - loss: 0.3079 -  
accuracy: 0.8843  
Epoch 7/20  
134/134 [=====] - 0s 1ms/step - loss: 0.2852 -  
accuracy: 0.8962  
Epoch 8/20  
134/134 [=====] - 0s 1ms/step - loss: 0.2673 -  
accuracy: 0.8976  
Epoch 9/20  
134/134 [=====] - 0s 1ms/step - loss: 0.2511 -  
accuracy: 0.9046  
Epoch 10/20  
134/134 [=====] - 0s 1ms/step - loss: 0.2394 -  
accuracy: 0.9085  
Epoch 11/20  
134/134 [=====] - 0s 1ms/step - loss: 0.2299 -  
accuracy: 0.9113  
Epoch 12/20  
134/134 [=====] - 0s 1ms/step - loss: 0.2202 -  
accuracy: 0.9162  
Epoch 13/20  
134/134 [=====] - 0s 1ms/step - loss: 0.2112 -  
accuracy: 0.9197  
Epoch 14/20  
134/134 [=====] - 0s 1ms/step - loss: 0.2049 -  
accuracy: 0.9246  
Epoch 15/20  
134/134 [=====] - 0s 1ms/step - loss: 0.2014 -  
accuracy: 0.9230  
Epoch 16/20
```

```

134/134 [=====] - 0s 1ms/step - loss: 0.1930 -
accuracy: 0.9279
Epoch 17/20
134/134 [=====] - 0s 1ms/step - loss: 0.1881 -
accuracy: 0.9293
Epoch 18/20
134/134 [=====] - 0s 1ms/step - loss: 0.1850 -
accuracy: 0.9342
Epoch 19/20
134/134 [=====] - 0s 1ms/step - loss: 0.1805 -
accuracy: 0.9323
Epoch 20/20
134/134 [=====] - 0s 1ms/step - loss: 0.1752 -
accuracy: 0.9370
<tensorflow.python.keras.callbacks.History at 0x7f1746587e80>

```

7) Quantify the model

We use our testing data to validate our model. This is how we determine the validity of our model (i.e. the ability to predict new and previously unseen data points)

```

model_loss, model_accuracy = model.evaluate(
    X_test_scaled, y_test_categorical, verbose=2)
print(
    f"Normal Neural Network - Loss: {model_loss}, Accuracy: {model_accuracy}")

```

```

45/45 - 0s - loss: 0.3214 - accuracy: 0.8929
Normal Neural Network - Loss: 0.32138586044311523, Accuracy:
0.892932116985321

```

8) Making Predictions with the new data

We use our trained model to make predictions using `model.predict`

```

predictions = model.predict_classes(X_test_scaled)

array([2, 0, 2, ..., 1, 1, 2])

y_test_categorical
array([[0., 0., 1.],
       [1., 0., 0.],
       [0., 0., 1.],
       ...,
       [0., 1., 0.],
       [0., 1., 0.],
       [0., 0., 1.]], dtype=float32)

true_labels = label_encoder.inverse_transform(predictions)
array(['data_scientist', 'data_analyst', 'data_scientist', ...,

```

```
'data_engineer', 'data_engineer', 'data_scientist'], dtype=object)
```

```
print(f"Predicted classes: {true_labels[:5]}")
```

```
print(f"Actual Labels: {list(y_test[:5])}")
```

```
Predicted classes: ['data_scientist' 'data_analyst' 'data_scientist'  
'data_scientist'  
'data_engineer']
```

```
Actual Labels: ['data_scientist', 'data_analyst', 'data_scientist',  
'data_analyst', 'data_engineer']
```